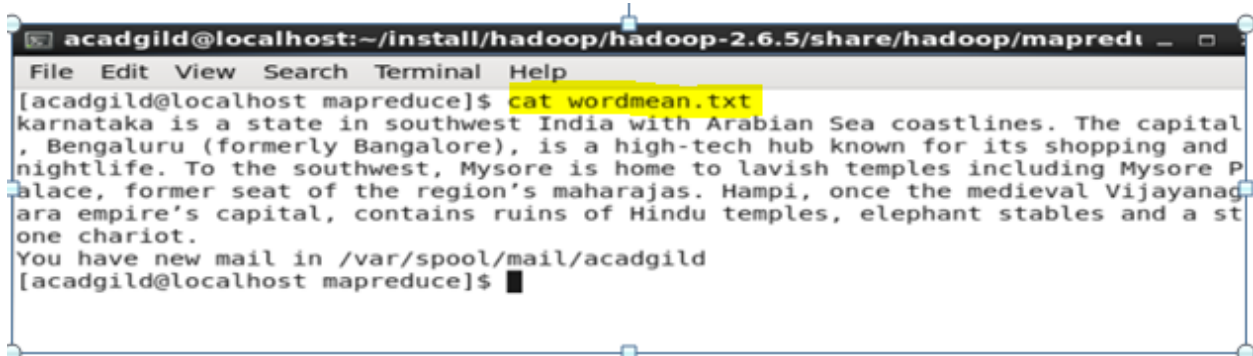


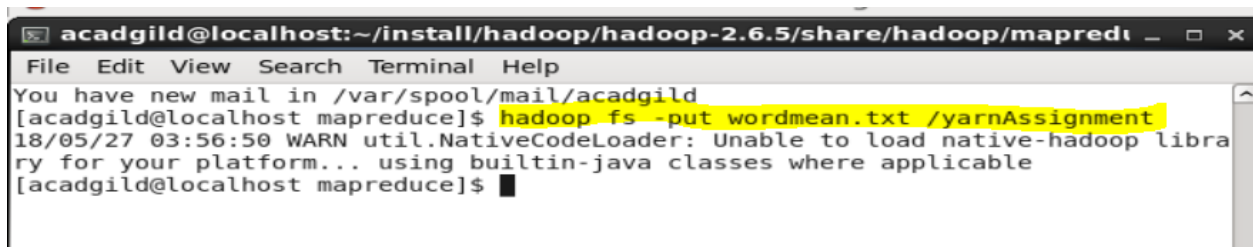
Task 1: Execute Word Mean program using hadoop mapreduce example

- a. Created wordmean.txt in local file system

A terminal window titled 'acadgild@localhost:~/install/hadoop/hadoop-2.6.5/share/hadoop/mapred...' with a menu bar (File, Edit, View, Search, Terminal, Help). The prompt is '[acadgild@localhost mapreduce]\$'. The command 'cat wordmean.txt' is entered and highlighted in yellow. The output shows text about Karnataka and a mail notification. The prompt returns to '[acadgild@localhost mapreduce]\$' after the output.

```
acadgild@localhost:~/install/hadoop/hadoop-2.6.5/share/hadoop/mapred...  
File Edit View Search Terminal Help  
[acadgild@localhost mapreduce]$ cat wordmean.txt  
karnataka is a state in southwest India with Arabian Sea coastlines. The capital  
, Bengaluru (formerly Bangalore), is a high-tech hub known for its shopping and  
nightlife. To the southwest, Mysore is home to lavish temples including Mysore P  
alace, former seat of the region's maharajas. Hampi, once the medieval Vijayanag  
ara empire's capital, contains ruins of Hindu temples, elephant stables and a st  
one chariot.  
You have new mail in /var/spool/mail/acadgild  
[acadgild@localhost mapreduce]$
```

- b. Copied wordmean.txt file into hadoop fs in /yarnAssignment directory

A terminal window with the same title and menu bar. The prompt is '[acadgild@localhost mapreduce]\$'. The command 'hadoop fs -put wordmean.txt /yarnAssignment' is entered and highlighted in yellow. The output shows a mail notification and a warning message. The prompt returns to '[acadgild@localhost mapreduce]\$' after the output.

```
acadgild@localhost:~/install/hadoop/hadoop-2.6.5/share/hadoop/mapred...  
File Edit View Search Terminal Help  
You have new mail in /var/spool/mail/acadgild  
[acadgild@localhost mapreduce]$ hadoop fs -put wordmean.txt /yarnAssignment  
18/05/27 03:56:50 WARN util.NativeCodeLoader: Unable to load native-hadoop libra  
ry for your platform... using builtin-java classes where applicable  
[acadgild@localhost mapreduce]$
```

- c. Executing **wordmean.txt** program using hadoop-mapreduce-example-2.9.0.jar and coping executed file into **/yarnAssignment/Output** Directory

```

[acadgild@localhost mapreduce]$ hadoop jar hadoop-mapreduce-examples-2.6.5.jar wordmean /yarnAssignment/wordmean.txt /yarnAssignment/output
18/05/27 04:08:29 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
18/05/27 04:08:31 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
18/05/27 04:08:36 INFO input.FileInputFormat: Total input paths to process : 1
18/05/27 04:08:36 INFO mapreduce.JobSubmitter: number of splits:1
18/05/27 04:08:37 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1527365639903_0001
18/05/27 04:08:39 INFO impl.YarnClientImpl: Submitted application application_1527365639903_0001
18/05/27 04:08:39 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1527365639903_0001/
18/05/27 04:08:39 INFO mapreduce.Job: Running job: job_1527365639903_0001
18/05/27 04:09:08 INFO mapreduce.Job: Job job_1527365639903_0001 running in uber mode : false
18/05/27 04:09:08 INFO mapreduce.Job: map 0% reduce 0%
18/05/27 04:09:24 INFO mapreduce.Job: map 100% reduce 0%
18/05/27 04:09:44 INFO mapreduce.Job: map 100% reduce 100%
18/05/27 04:09:45 INFO mapreduce.Job: Job job_1527365639903_0001 completed successfully
18/05/27 04:09:46 INFO mapreduce.Job: Counters: 49
  File System Counters
    FILE: Number of bytes read=39
    FILE: Number of bytes written=215447
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=531
    HDFS: Number of bytes written=20
    HDFS: Number of read operations=6
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=1
    Launched reduce tasks=1
    Data-local map tasks=1
    Total time spent by all maps in occupied slots (ms)=12652
    Total time spent by all reduces in occupied slots (ms)=17535
    Total time spent by all map tasks (ms)=12652
    Total time spent by all reduce tasks (ms)=17535
    Total vcore-milliseconds taken by all map tasks=12652
    Total vcore-milliseconds taken by all reduce tasks=17535
    Total megabyte-milliseconds taken by all map tasks=12955648
    Total megabyte-milliseconds taken by all reduce tasks=17955840
Map-Reduce Framework

```

d. The mean value of **wordmean.txt** file.

```

-----
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=417
File Output Format Counters
  Bytes Written=20
The mean is: 5.661290322580645
-----

```

e. localhost:50070 – UI representation

/yarnAssignment/output						Go!
Permission	Owner	Group	Size	Replication	Block Size	Name
-rw-r--r--	acadgild	supergroup	0 B	1	128 MB	SUCCESS
-rw-r--r--	acadgild	supergroup	20 B	1	128 MB	part-r-00000

Task 2: Execute Word Median program using hadoop mapreduce example

- a. Created **wordmedian.txt** in local file system

```
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost mapreduce]$ vi wordmedian.txt
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost mapreduce]$ cat wordmedian.txt
India is a vast South Asian country with diverse terrain – from Himalayan peaks to Indian Ocean coastline – and history reaching back 5 millennia. In the north, Mughal Empire landmarks include Delhi's Red Fort complex and massive Jama Masjid mosque, plus Agra's iconic Taj Mahal mausoleum. Pilgrims bathe in the Ganges in Varanasi, and Rishikesh is a yoga centre and base for Himalayan trekking
[acadgild@localhost mapreduce]$
```

- b. Copied **wordmedian.txt** file into hadoop fs in /yarnAssignment directory

- c. Executing **wordmean.txt** program using hadoop-mapreduce-example-2.9.0.jar and copying executed file into /yarnAssignment/**OutputMedian** Directory

```
[acadgild@localhost mapreduce]$ hadoop fs -put wordmedian.txt /yarnAssignment
18/05/27 04:16:57 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost mapreduce]$ hadoop jar hadoop-mapreduce-examples-2.6.5.jar wordmedian /yarnAssignment/wordmedian.txt /yarnAssignment/outputMedian
18/05/27 04:17:42 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
18/05/27 04:17:45 INFO client.RMPProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
18/05/27 04:17:48 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
18/05/27 04:17:49 INFO input.FileInputFormat: Total input paths to process : 1
18/05/27 04:17:49 INFO mapreduce.JobSubmitter: number of splits:1
18/05/27 04:17:50 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1527365639903_0002
18/05/27 04:17:51 INFO impl.YarnClientImpl: Submitted application application_1527365639903_0002
18/05/27 04:17:51 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1527365639903_0002/
18/05/27 04:17:51 INFO mapreduce.Job: Running job: job_1527365639903_0002
18/05/27 04:18:13 INFO mapreduce.Job: Job job_1527365639903_0002 running in uber mode : false
18/05/27 04:18:13 INFO mapreduce.Job: map 0% reduce 0%
18/05/27 04:18:26 INFO mapreduce.Job: map 100% reduce 0%
18/05/27 04:18:44 INFO mapreduce.Job: map 100% reduce 100%
18/05/27 04:18:45 INFO mapreduce.Job: Job job_1527365639903_0002 completed successfully
18/05/27 04:18:45 INFO mapreduce.Job: Counters: 49

File System Counters
  FILE: Number of bytes read=106
  FILE: Number of bytes written=215333
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=519
  HDFS: Number of bytes written=41
  HDFS: Number of read operations=6
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=2

Job Counters
  Launched map tasks=1
  Launched reduce tasks=1
  Data-local map tasks=1
  Total time spent by all maps in occupied slots (ms)=11218
  Total time spent by all reduces in occupied slots (ms)=13942
  Total time spent by all map tasks (ms)=11218
  Total time spent by all reduce tasks (ms)=13942
  Total vcore-milliseconds taken by all map tasks=11218
  Total vcore-milliseconds taken by all reduce tasks=13942
  Total megabyte-milliseconds taken by all map tasks=11487232
  Total megabyte-milliseconds taken by all reduce tasks=14276608
```

- d. The median value of wordmedian.txt file.

```

Total megabyte-milliseconds taken by all reduce tasks=14271
Map-Reduce Framework
  Map input records=1
  Map output records=65
  Map output bytes=520
  Map output materialized bytes=106
  Input split bytes=116
  Combine input records=65
  Combine output records=10
  Reduce input groups=10
  Reduce shuffle bytes=106
  Reduce input records=10
  Reduce output records=10
  Spilled Records=20
  Shuffled Maps =1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=266
  CPU time spent (ms)=3130
  Physical memory (bytes) snapshot=291549184
  Virtual memory (bytes) snapshot=4118224896
  Total committed heap usage (bytes)=170004480
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=403
File Output Format Counters
  Bytes Written=41
The median is: 5

```

e. localhost:50070 – UI representation

/yarnAssignment/outputMedian						Go!
Permission	Owner	Group	Size	Replication	Block Size	Name
-rw-r--r--	acadgild	supergroup	0 B	1	128 MB	_SUCCESS
-rw-r--r--	acadgild	supergroup	41 B	1	128 MB	part-r-00000

Task 3: Execute Word Standard Deviation program using hadoop mapreduce example

- Created **wordStandardDeviation.txt** in local file system
- Copied **wordStandardDeviation.txt** file into hadoop fs in /yarnAssignment directory
- Executing wordmean.txt program using hadoop-mapreduce-example-2.9.0.jar and coping executed file into /yarnAssignment/**OutputStandardDeviation** Directory

```

You have new mail in /var/spool/mail/acadgild
[acadgild@localhost mapreduce]$ vi wordStandardDeviation.txt
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost mapreduce]$ cat wordStandardDeviation.txt
Bengaluru (also called Bangalore) is the capital of India's southern Karnataka state. The center of India's high-tech indus
try, the city is also known for its parks and nightlife. By Cubbon Park, Vidhana Soudha is a Neo-Dravidian legislative buil
ding. Former royal residences include 19th-century Bangalore Palace, modeled after England's Windsor Castle, and Tipu Sulta
n's Summer Palace, an 18th-century teak structure.
[acadgild@localhost mapreduce]$ hadoop fs -put wordStandardDeviation.txt /yarnAssignment
18/05/27 04:24:23 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java
Classes where applicable
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost mapreduce]$ hadoop jar hadoop-mapreduce-examples-2.6.5.jar wordStandardDeviation /yarnAssignment/wordSt
andardDeviation.txt /yarnAssignment/outputStandardDeviation

```

- The mean value of **wordStandardDeviation.txt** file.

```

    total megabyte-milliseconds taken by all reduce tasks=16420864
Map-Reduce Framework
  Map input records=1
  Map output records=177
  Map output bytes=2596
  Map output materialized bytes=56
  Input split bytes=127
  Combine input records=177
  Combine output records=3
  Reduce input groups=3
  Reduce shuffle bytes=56
  Reduce input records=3
  Reduce output records=3
  Spilled Records=6
  Shuffled Maps =1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=297
  CPU time spent (ms)=3290
  Physical memory (bytes) snapshot=293990400
  Virtual memory (bytes) snapshot=4118220800
  Total committed heap usage (bytes)=170004480
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=424
File Output Format Counters
  Bytes Written=32
The standard deviation is: 2.9463263294166797

```

e. localhost:50070 – UI representation

/yarnAssignment/outputStandardDeviation						Go!
Permission	Owner	Group	Size	Replication	Block Size	Name
-rw-r--r--	acadgild	supergroup	0 B	1	128 MB	_SUCCESS
-rw-r--r--	acadgild	supergroup	32 B	1	128 MB	part-r-00000

The Final output files of all 3 tasks.

Permission	Owner	Group	Size	Replication	Block Size	Name
drwxr-xr-x	acadgild	supergroup	0 B	0	0 B	output
drwxr-xr-x	acadgild	supergroup	0 B	0	0 B	outputMedian
drwxr-xr-x	acadgild	supergroup	0 B	0	0 B	outputStandardDeviation
-rw-r--r--	acadgild	supergroup	424 B	1	128 MB	wordStandardDeviation.txt
-rw-r--r--	acadgild	supergroup	417 B	1	128 MB	wordmean.txt
-rw-r--r--	acadgild	supergroup	403 B	1	128 MB	wordmedian.txt